

# Machine Learning – Homework 2

Andrea Gasparini  
1813486

December 2020

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Dataset</b>	<b>2</b>
<b>3</b>	<b>Preprocessing</b>	<b>3</b>
3.1	Data augmentation . . . . .	4
<b>4</b>	<b>Convolutional Neural Networks</b>	<b>4</b>
4.1	AlexNet . . . . .	4
4.2	Custom network . . . . .	6
<b>5</b>	<b>Evaluation</b>	<b>6</b>
5.1	AlexNet . . . . .	7
5.2	Custom network . . . . .	8
<b>6</b>	<b>Conclusions</b>	<b>10</b>

# 1 Introduction

The goal of this Homework was to solve an image classification problem with objects typically available in a home environment.

## 2 Dataset

The dataset contains **8221** samples and it's based on a 8 classes subset of the [RoboCup@Home-Objects dataset](#). For each class we have a folder containing all the related samples.

- *accent plate*
- *breadsticks*
- *cereals box*
- *cocoa drink bottle*
- *nitrile gloves*
- *paper bag*
- *pears*
- *plastic fork*




Figure 1: Distribution of the dataset's 8 classes

As we can see in Figure 1, the dataset is quite balanced, but it should be noted that it also contains outliers – i.e. misclassified images – meaning that our model should adapt accordingly.



Figure 2: 10 random samples from the dataset

### 3 Preprocessing

The first thing I did is the dataset splitting in two subsets for training and validation, respectively 70% and 30% of the original one. The splitting has been made at filesystem level, obtaining two directories (*train* and *val*) with the same structure of the original one (i.e. with 8 subdirectories, one for each class).

### 3.1 Data augmentation

Data augmentation includes techniques used to generate *new* training samples from the original ones by applying random jitters and transformations, but at the same time ensuring that the class labels of the data are not changed. These kind of techniques are useful to reduce overfitting and they usually provide an increase in testing accuracy, at the expense of a slight worsening in training accuracy. In our case I obtained an increase in the performance by 5 to 10% applying flipping, rotations and zooming to the training set of images.

## 4 Convolutional Neural Networks

All the tested solutions for this problem are based on a **Convolutional Neural Network** (CNN), commonly composed by convolutional layers, pooling layers and fully-connected layers. The **convolutional layers** apply a convolution operation to the input images, the **pooling layer** consequently perform a down-sampling through a non linear function and the **fully connected layers** finally classify in one of the 8 classes.

### 4.1 AlexNet

The first tested network is the implementation of AlexNet, a popular CNN, which has eight layers using the non-saturating **ReLU** activation function:

- five convolutional layers, some of them followed by a max-pooling layer
- three fully connected layers, each one followed by a Dropout layer

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 120, 60, 96)	34944
max_pooling2d (MaxPooling2D)	(None, 60, 30, 96)	0
batch_normalization (Batch Normalization)	(None, 60, 30, 96)	384
conv2d_1 (Conv2D)	(None, 50, 20, 256)	2973952
max_pooling2d_1 (MaxPooling2D)	(None, 25, 10, 256)	0
batch_normalization_1 (Batch Normalization)	(None, 25, 10, 256)	1024
conv2d_2 (Conv2D)	(None, 23, 8, 384)	885120
batch_normalization_2 (Batch Normalization)	(None, 23, 8, 384)	1536
conv2d_3 (Conv2D)	(None, 21, 6, 384)	1327488
batch_normalization_3 (Batch Normalization)	(None, 21, 6, 384)	1536
conv2d_4 (Conv2D)	(None, 19, 4, 256)	884992
max_pooling2d_2 (MaxPooling2D)	(None, 9, 2, 256)	0
batch_normalization_4 (Batch Normalization)	(None, 9, 2, 256)	1024
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 4096)	18878464
dropout (Dropout)	(None, 4096)	0
batch_normalization_5 (Batch Normalization)	(None, 4096)	16384
dense_1 (Dense)	(None, 4096)	16781312
dropout_1 (Dropout)	(None, 4096)	0
batch_normalization_6 (Batch Normalization)	(None, 4096)	16384
dense_2 (Dense)	(None, 1000)	4097000
dropout_2 (Dropout)	(None, 1000)	0
batch_normalization_7 (Batch Normalization)	(None, 1000)	4000
dense_3 (Dense)	(None, 8)	8008
Total params: 45,913,552		
Trainable params: 45,892,416		
Non-trainable params: 21,136		

Figure 3: Summary of AlexNet

## 4.2 Custom network

After the evaluation of AlexNet, the main analysis has been made with a custom network, in which it was possible to test different layers and hyper-parameters. For example adding dropout layers or changing the dropout rate. The optimal configuration obtained is showed in Figure 4, all of them using the ReLU activation function except for the last one with `softmax`

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 250, 250, 16)	448
max_pooling2d (MaxPooling2D)	(None, 125, 125, 16)	0
conv2d_1 (Conv2D)	(None, 125, 125, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 32)	0
conv2d_2 (Conv2D)	(None, 62, 62, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 64)	0
dropout (Dropout)	(None, 31, 31, 64)	0
flatten (Flatten)	(None, 61504)	0
dense (Dense)	(None, 128)	7872640
dense_1 (Dense)	(None, 8)	1032
Total params: 7,897,256		
Trainable params: 7,897,256		
Non-trainable params: 0		

Figure 4: Summary of the custom network

## 5 Evaluation

After some test with different epochs values I fixed it at 20 for the final evaluation, which was the optimal value for performance and training time. To compare the results, in this section are available the accuracy and loss curves and also the classification report for both AlexNet and the custom network. The accuracy curve shows how the network accuracy increase over every epoch, while the loss curve shows how the loss decrease (hopefully).

## 5.1 AlexNet

As we can see in Figure 5, while the accuracy of the training increase and the loss decrease for every epoch, the validation results are very poor. The gap between the training curve and the validation curve also measure how much our model is prone to overfitting (the model adapt too much on the data in the training while performing poor on the overall data available).




Figure 5: AlexNet accuracy and loss plots

	<b>Precision</b>	<b>Recall</b>	<b>F1-score</b>	<b>Support</b>
<i>breadsticks</i>	0.685	0.361	0.473	313
<i>cereals box</i>	0.853	0.625	0.721	344
<i>cocoa drink bottle</i>	0.653	0.446	0.530	296
<i>nitrile gloves</i>	0.651	0.825	0.728	309
<i>pears</i>	0.399	0.915	0.556	306
<i>accent plate</i>	0.813	0.742	0.776	322
<i>paper bag</i>	0.772	0.392	0.520	355
<i>plastic fork</i>	0.572	0.721	0.638	226
<hr/>				
<b>Accuracy</b>			0.622	2471
<b>Macro avg</b>	0.675	0.628	0.618	2471
<b>Weighted avg</b>	0.684	0.622	0.618	2471

Table 1: Classification report of AlexNet

## 5.2 Custom network

The following results for the custom network are based on the configuration with a dropout rate of 0.4, which is the optimal one compared to other values. The performance has been increased and regularized compared to the AlexNet ones that were unbalanced in precision and recall (e.g. the *pears* class had a low precision of 0.399 and a high recall of 0.915).






Figure 6: Custom network accuracy and loss plots

	Precision	Recall	F1-score	Support
<i>breadsticks</i>	0.814	0.812	0.813	313
<i>cereals box</i>	0.816	0.762	0.788	344
<i>cocoa drink bottle</i>	0.568	0.622	0.594	296
<i>nitrile gloves</i>	0.871	0.702	0.778	309
<i>pears</i>	0.813	0.768	0.790	306
<i>accent plate</i>	0.863	0.786	0.823	322
<i>paper bag</i>	0.562	0.769	0.649	355
<i>plastic fork</i>	0.751	0.655	0.700	226
<b>Accuracy</b>			0.739	2471
<b>Macro avg</b>	0.757	0.734	0.742	2471
<b>Weighted avg</b>	0.756	0.739	0.743	2471

Table 2: Classification report of the custom network

## 6 Conclusions

The performances have been unexpectedly better with the final custom network compared to AlexNet, as highlighted by the accuracy and loss curves. The more likely reason is that a network like this needs to have a larger dataset and a bigger number of epochs to obtain good results. In fact I got better results with 100 epochs, but that required more than 1 hour of training. In conclusion, assuming that there isn't overfitting, we can say that the proposed custom network is a good fit for this problem. To obtain the best performances it's probably necessary to use transfer learning from a pre-trained network on a better dataset.